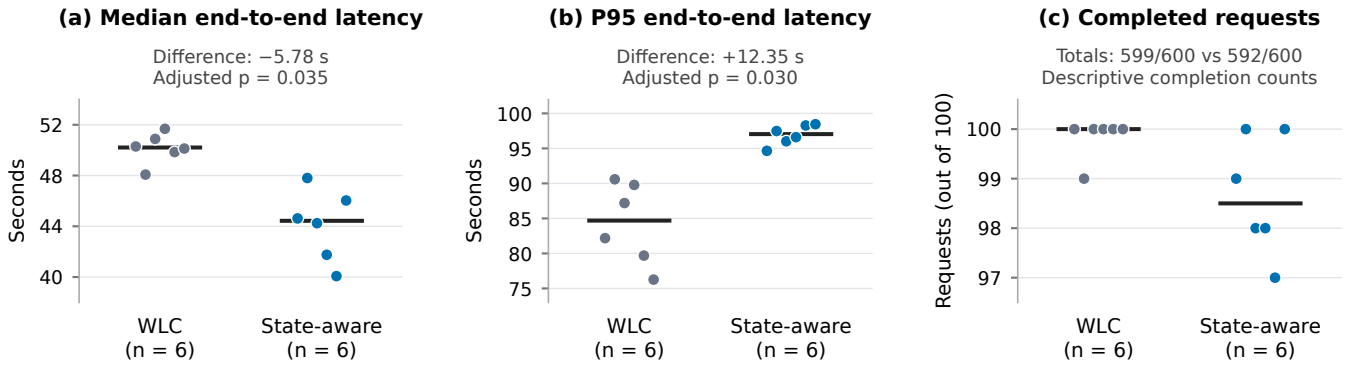
(a) Median end-to-end latency
Difference: −5.78 s
Adjusted p = 0.035
Seconds
52
48
44
40
WLC
(n = 6)
State-aware
(n = 6)
(b) P95 end-to-end latency
Difference: +12.35 s
Adjusted p = 0.030
Seconds
100
95
90
85
80
75
WLC
(n = 6)
State-aware
(n = 6)
(c) Completed requests
Totals: 599/600 vs 592/600
Descriptive completion counts
Requests (out of 100)
100
99
98
97
WLC
(n = 6)
State-aware
(n = 6)
ShareGPT c12 • Each point: one run • Black bars: medians across runs